

SQL Server: Query Plan Analysis

Module 4: Common Operators

Joe Sack

Joe@SQLskills.com





Module Introduction

- **This module will review the more common operators that you may expect to see in a query execution plan**
- **Don't believe statements that specific operators or behaviors translate to an absolute problem**
 - I'll call out things to watch for in the next module, but again, see them as areas of investigation


Table and Index Scans

- **Table scan**
 - Indicating a retrieval of ALL rows from a table without a clustered index
- **Clustered Index scan**
 - Indicating a retrieval of ALL rows from a table with a clustered index
- **Columnstore Index Scan**
 - New as of SQL Server 2012, columnar storage index



Index Seeks

- **Clustered Index Seek**
 - Retrieving rows based on a SEEK predicate from clustered index
- **Nonclustered Index Seek**
 - Same, but from a nonclustered index




Lookups

- **Key Lookup**

- Bookmark lookup on table with clustered index (always via Nested Loop)

- **RID Lookup**

- Is just a bookmark lookup to a heap (using the RID)
 - Just like with Key Lookups, you'll only see this with Nested Loop Joins



Join Considerations

- Beware of advice telling you that specific join types (or operators, for that matter) are “good” or “bad”
- Key factors:
 - Tables to be joined, order of joins, algorithm used, cardinality
- Join hints and/or forcing order = red flag
 - Generally, “edge” cases or extreme tuning scenarios require their use
 - Otherwise, ask questions and find out why this is happening

Outer/Inner Terminology

- **When it comes to joins, you may hear the “outer” vs. “inner” table terminology**
 - But it is NOT related to the order you write them in your query, unless you’re forcing it
 - Outer = top = left
 - Nested Loop: for each row in outer, find all rows in inner
 - Merge Join: inner/outer – not as important (will discuss why)
 - Hash Join: outer table is the “build” hash table
 - Inner = bottom = right
 - Hash Match: inner table is probe table


Nested Loop

- **Supports:**

- Inner join, left outer join, left semi join, left anti semi join, cross join, cross apply, outer apply


- **Algorithm:**

- For one row in the outer (top) table, find matching rows in the inner (bottom) table and return them
- After no matching rows on the inner table are found, retrieve the next row from the outer (top) table and repeat until end of outer (top) table rows





Merge Join

- **Logical operations:**
 - Requires one equijoin predicate
 - Except for full outer join transformation in many-to-many scenario
- **Joins two inputs (sorted on joining columns)**
 - Pre-existing sorting (via index) is ideal, but sorts can be automatically added (noteworthy if you see this)
- **Algorithm:**
 - Retrieve row from outer and inner tables
 - If a match: return the row
 - If no match: get a new row from the smaller input and iterate



Hash Match Join

- Requires one equijoin predicate
- Joins can be on unsorted inputs
- **Algorithm:**
 - Build a hash table (hash buckets) via computed hash key values for each row of the “build” input (top/outer table)
 - For each probe row (bottom/inner table), compute a hash key value and evaluate for matches in the “build” hash table (buckets)
 - Output matches (or output based on logical operation)
- **Build Hash**
 - Build of a batch hash table for a columnstore index




Filter

- **Predicates can be evaluated within operators that read data from table/indexes**
- **Query Optimizer aims (when possible) to “push” filter down the tree (leaf level) to reduce rows moved**
- **Sometimes a predicate cannot be pushed and you’ll see a Filter operator instead**
- **When you see these, take note of where they are happening**
 - Late in the data flow can translate to higher overhead as the operators pull data




Stream Aggregate

- **Groups rows by one or more columns**
 - Calculates one or more aggregate expressions
 - Does this one group at a time
 - Can do:
 - Scalar aggregates (no GROUP BY) e.g. SUM
 - Group aggregates (have GROUP BY)
 - Requires ordered (sorted) input for grouping columns
 - If unordered then Query Optimizer can add a Sort operator
- **Non blocking (streams)**
- **Minimal memory**




Hash Match (Aggregate)

- **Several aspects of hash join apply to hash match**
 - Requires memory for hash table
 - Unsorted inputs are okay
 - Is stop-and-go on the build table
- **For hash match, hash table generated for groupings of rows**
 - Hash table values based on grouping columns
 - 1) Generate hash
 - 2) Check for existing row in hash table
 - 3) Generate row if no match or update matching row



Sort

- **As named: orders rows received from input**
 - For example, due to ORDER BY in query
- **Variation is “Distinct Sort” to remove duplicates**



Spools

■ Eager Spool

- Takes entire input, placing row(s) in hidden tempdb work file
- When spool's parent operator asks for first row, spool grabs all rows from the input and stores them in tempdb

■ Lazy Spool

- When the lazy spool's parent requests a row, the spool grabs the row from the input operator and stores it in the tempdb spool table
 - It does not retrieve all rows like the eager spool
- Lazy spools are non-blocking




Eager Spool and “Halloween Protection”

- **Identified by IBM researchers back in 1976**
 - Performing an update involved, conceptually a read cursor and a write cursor
 - The write cursor was updating the read cursor, causing a moving target of re-updates to the same set of rows
- **Eager spools, when needed, prevent the write cursor from impacting the read cursor**
 - Example of when its not needed?
 - If you're NOT updating the index key itself (causing movement that could make rows scanned > 1)
 - Other blocking operators already in use (such as Sort)

Constant Scan

- **Introduces ≥ 1 constant rows that can be built upon by parent operators**
 - You can see this in data modification plans as well as SELECT plans
- **SQL Server 2005: seen with partitioning as well, defining applicable partition numbers (driven by predicates)**
 - Nested Loop operator joining Constant Scan (outer) and partitioned table (inner)




Assert

- **Verifies existence of a specific condition**
 - Check constraints
 - Referential integrity
 - Enforce return of one-row for scalar sub-query




Compute Scalar

- Evaluates an expression, producing a scalar value (e.g. GETDATE() value)
- Estimated CPU cost is often low – for example 0.001 for an estimated 9,615 rows passing through an UPPER function
 - May be a placeholder definition of an expression calculated in another operator – for example...
 - Compute Scalar
 - [Expr1003] = Scalar Operator(upper([Credit].[dbo].[member].[lastname]))
 - Sort
 - [Credit].[dbo].[member].member_no, [Credit].[dbo].[member].curr_balance, Expr1003



Identifying Parallelism in the Plan

- Parallelism operators (Distribution/Repartition/Gather)
- You'll also see the parallelism icon in the graphic for operators that can run in parallel or serial modes



- If looking at XML Showplan you'll see RelOp:
 - Parallel="true"

Exchange Operators

- **Distribute Streams**




- Takes one input and produces multiple output data streams

- **Repartition Streams**

- Takes in multiple streams and then produces multiple streams
 - Can be used in conjunction with bitmap filter to reduce rows in output
 - May “rebalance” thread skew


- **Gather Streams**

- Operator takes in multiple streams and produces a single stream




Bitmap

- **Performs bitmap filtering for parallel plans with hash or merge joins**
 - Optimization that eliminates rows with key values that wouldn't produce join records
 - Reduces rows being passed to parent operators




Merge Interval

- Merges multiple potentially overlapping intervals used in predicates
- Preceded by compute scalar and constant scan
- Goal is to minimize redundant scans



Concatenation

- Scans multiple inputs, returning each row
- Seen via UNION ALL



Segment and Sequence Project

- **Segment divides input rows into related segments based on specific columns**
 - Seen with windowing functions
 - Columns used for segmenting shown in argument properties
 - Extra segment column created which tracks if value has changed from the previous row
- **Sequence Project adds columns to perform computations on an ordered set**
 - Outputs one segment at a time

